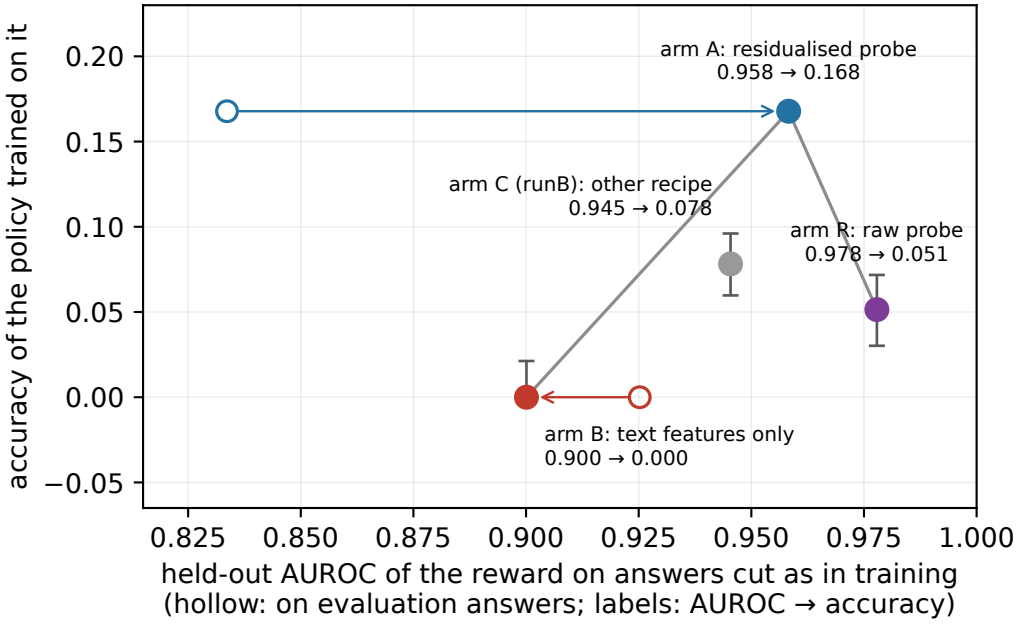
accuracy of the policy trained on it
0.20
0.15
0.10
0.05
0.00
−0.05
0.825
0.850
0.875
0.900
0.925
0.950
0.975
1.000
held-out AUROC of the reward on answers cut as in training
(hollow: on evaluation answers; labels: AUROC → accuracy)
arm A: residualised probe
0.958 → 0.168
arm C (runB): other recipe
0.945 → 0.078
arm R: raw probe
0.978 → 0.051
arm B: text features only
0.900 → 0.000